

# Citibike Case Study By Mahwish Khalid

## Background, Context, and Objective

Client: Mayor of New York City, Bill de Blasio Objective: Help the mayor get a better understanding of citibike ridership by creating an operating report for 2017. Ask: 1. Top 5 stations with the most starts (showing # of starts) 2. Trip duration by user type 3. Most popular trips based on start station and stop station) 4. Rider performance by Gender and Age based on avg trip distance (station to station), median speed (distance traveled / trip duration) 5. What is the busiest bike in NYC in 2017? How many times was it used? How many minutes was it in use? Note: A model that can predict how long a trip will take given a starting point and destination. Solution:- The dataset is massive load the dataset and import the libraries before loading te dataset.

## QUESTION 1

Question 1: Top 5 Stations Let's check if there's any noise or cleanup which needs to be done before creating the chart. Any missing values? Mostly for Birth year and a few for User Type. We can ignore these for now and deal with them later. Let's get the data in the right format Trip Duration - Int Start Time - DateTime Stop Time - DateTime Start Station ID - Categorical Start Station Name - Categorical User Type - Categorical Birth Year - Ordinal Gender - Categorical Deal with trips which lasted less than 1.5 minute (90 seconds). If so, in the ideal world, we should not include this start, we may double count. If a bike is broken, a user will dock it again within a minute or two and pick-up another one.

## Import Libraries

```
In [57]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas import *
import seaborn as sns
from matplotlib import rcParams
import datetime as dt
```

## Import csv file

```
In [58]: df=pd.read_csv("C:/Users/Mahwish/Desktop/mc master university/case study IBM/citibike.csv")
```

## Basic Summary,dimension and Structure of Dataset

### 1)First Few Rows of Dataset

```
In [59]: print(df.head())
```

	Trip Duration	Start Time	Stop Time	Start Station ID \
0	680	2017-01-01 00:00:21	2017-01-01 00:11:41	3226
1	1282	2017-01-01 00:00:45	2017-01-01 00:22:08	3263
2	648	2017-01-01 00:00:57	2017-01-01 00:11:46	3143
3	631	2017-01-01 00:01:10	2017-01-01 00:11:42	3143
4	621	2017-01-01 00:01:25	2017-01-01 00:11:47	3143

	Start Station Name	Start Station Latitude \
0	W 82 St & Central Park West	40.782750
1	Cooper Square & E 7 St	40.729236
2	5 Ave & E 78 St	40.776829
3	5 Ave & E 78 St	40.776829
4	5 Ave & E 78 St	40.776829

	Start Station Longitude	End Station ID	End Station Name \
0	-73.971370	3165	Central Park West & W 72 St
1	-73.990868	498	Broadway & W 32 St
2	-73.963888	3152	3 Ave & E 71 St
3	-73.963888	3152	3 Ave & E 71 St
4	-73.963888	3152	3 Ave & E 71 St

	End Station Latitude	End Station Longitude	Bike ID	User Type \
0	40.775794	-73.976206	25542	Subscriber
1	40.748549	-73.988084	21136	Subscriber
2	40.768737	-73.961199	18147	Customer
3	40.768737	-73.961199	21211	Customer
4	40.768737	-73.961199	26819	Customer

	Birth Year	Gender
0	1965.0	2
1	1987.0	2
2	NaN	0
3	NaN	0
4	NaN	0

## 2) Dimension of dataset

```
In [60]: df.shape
```

```
Out[60]: (726676, 15)
```

### 3)Data type of each column

```
In [61]: print(df.dtypes)
```

```
Trip Duration          int64
Start Time             object
Stop Time             object
Start Station ID       int64
Start Station Name     object
Start Station Latitude float64
Start Station Longitude float64
End Station ID         int64
End Station Name      object
End Station Latitude   float64
End Station Longitude  float64
Bike ID               int64
User Type             object
Birth Year            float64
Gender               int64
dtype: object
```

#### ***So Dataset contains categorical Features***

```
In [62]: categorical = df.dtypes[df.dtypes == "object"].index
```

```
In [63]: print(categorical)
```

```
Index(['Start Time', 'Stop Time', 'Start Station Name', 'End Station Name',
      'User Type'],
      dtype='object')
```

**Categorical features must be transformed into numerical features to be useful in most types of analysis.**

```
In [64]: df["Start Time"] = pd.to_datetime(df["Start Time"] )  
df["Stop Time"] = pd.to_datetime(df["Stop Time"] )
```

```
In [65]: df["Start Station Name"] = df["Start Station Name"].astype('category')  
df["End Station Name"] = df["End Station Name"].astype('category')  
df["User Type"] = df["User Type"].astype('category')
```

#### **4)Statistical Summary of data**

```
In [66]: print(df.describe())
```

	Trip Duration	Start Station ID	Start Station Latitude \
count	7.266760e+05	726676.000000	726676.000000
mean	7.778989e+02	1222.917630	40.737372
std	1.124683e+04	1277.955252	0.072596
min	6.100000e+01	72.000000	0.000000
25%	3.310000e+02	358.000000	40.720874
50%	5.260000e+02	482.000000	40.739355
75%	8.600000e+02	3092.000000	40.755103
max	5.325688e+06	3446.000000	40.804213

	Start Station Longitude	End Station ID	End Station Latitude \
count	726676.000000	726676.000000	726676.000000
mean	-73.984795	1197.252902	40.737077
std	0.123776	1266.085070	0.072474
min	-74.031372	72.000000	0.000000
25%	-73.995299	356.000000	40.720828
50%	-73.987167	479.000000	40.739323
75%	-73.976682	3078.000000	40.755003
max	0.000000	3447.000000	40.804213

	End Station Longitude	Bike ID	Birth Year	Gender
count	726676.000000	726676.000000	697600.000000	726676.000000
mean	-73.985133	21713.053902	1977.122481	1.166728
std	0.123782	4199.313576	11.925020	0.475971
min	-74.033459	14529.000000	1885.000000	0.000000
25%	-73.995960	17859.000000	1969.000000	1.000000
50%	-73.987586	21295.000000	1979.000000	1.000000
75%	-73.976806	25803.000000	1987.000000	1.000000
max	0.000000	27325.000000	2000.000000	2.000000

## Handling Missing Values ¶

### 1)Find Exact Number of Missing Values

```
In [67]: pd.isnull(obj=df).values.ravel().sum()
```

```
Out[67]: 32269
```

## 2) Total Number of Missing values in Column

```
In [68]: df.isnull().sum(axis=0)
```

```
Out[68]: Trip Duration          0
         Start Time            0
         Stop Time             0
         Start Station ID      0
         Start Station Name     0
         Start Station Latitude 0
         Start Station Longitude 0
         End Station ID        0
         End Station Name       0
         End Station Latitude   0
         End Station Longitude  0
         Bike ID               0
         User Type             3193
         Birth Year            29076
         Gender                0
         dtype: int64
```

***Mostly for Birth year and a few for User Type. So it safe to remove NA values of User Type and other unnecessary data.***

## 3) Deleting Unnecessary Data

```
In [69]: df = df.dropna(subset=['User Type'])
```

```
In [70]: df.isnull().sum(axis=0)
```

```
Out[70]: Trip Duration          0
         Start Time            0
         Stop Time             0
         Start Station ID      0
         Start Station Name     0
         Start Station Latitude 0
         Start Station Longitude 0
         End Station ID        0
         End Station Name       0
         End Station Latitude   0
         End Station Longitude  0
         Bike ID               0
         User Type              0
         Birth Year             29071
         Gender                 0
         dtype: int64
```

***Citi Bike riders often come across damage or broken bikes. Let's drop any trips where a trip lasted less than 90 seconds ,As we can see from statistical summary minimum duration is 61sec.Also we drop double counts inStart Station Ltitude and End Station Latitude***

```
In [71]: df = df.drop(df.index[(df['Trip Duration'] < 90) &
                               (df['Start Station Latitude'] == df['End Station Latitude'])])
```

```
In [72]: df.shape
```

```
Out[72]: (722528, 15)
```

## Question 1)Top 5 stations with the most starts (showing # of starts)

Data for Top 5 Stations visual



```
In [73]: Top5_Stations=df['Start Station Name'].value_counts().head().index  
Number_of_Start=df['Start Station Name'].value_counts().head().values  
Top5_Stations = Top5_Stations.astype('object')
```

```
In [74]: Top5_Stations
```

```
Out[74]: Index(['Pershing Square North', 'W 21 St & 6 Ave', 'E 17 St & Broadway',  
              'Broadway & E 22 St', '8 Ave & W 33 St'],  
              dtype='object')
```

```
In [75]: Number_of_Start
```

```
Out[75]: array([8760, 5435, 5099, 4900, 4789], dtype=int64)
```

```
In [ ]:
```

```
In [76]: import seaborn as sns
import matplotlib.pyplot as plt
% matplotlib inline

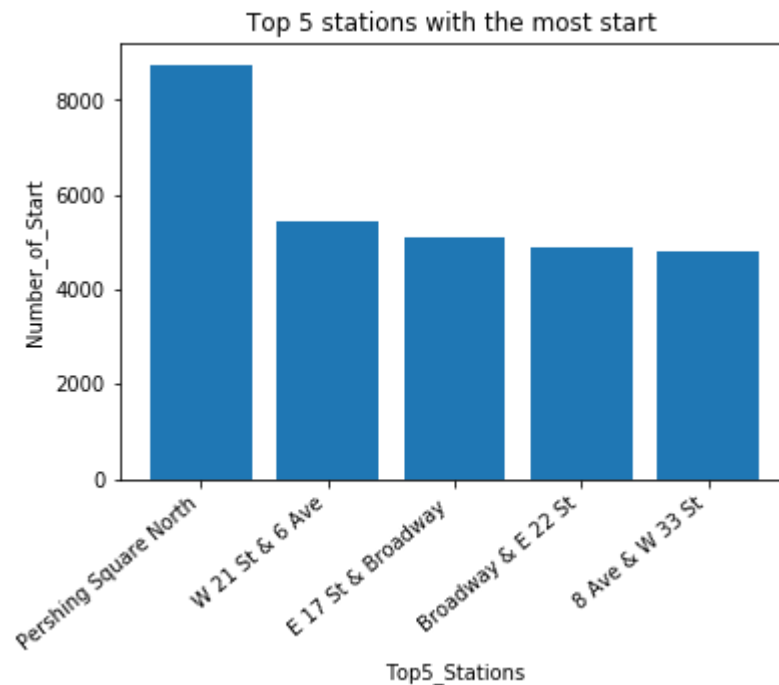
# Choose the names of the bars
Top5_Stations = ('Pershing Square North', 'W 21 St & 6 Ave', 'E 17 St & Broadway ', 'Broadway & E 22 St' , '8 A
ve & W 33 St')
y_pos = np.arange(len(Top5_Stations))

# Create bars
ax=plt.bar(y_pos,Number_of_Start)

# Create names on the x-axis
plt.xticks(y_pos,Top5_Stations , color='black')
plt.yticks(color='black')

# Create names on the x-axis any Y axis
plt.xticks(y_pos, Top5_Stations)
Title = "Top 5 stations with the most start "
Xl = "Top5_Stations"
Yl = "Number_of_Start"
plt.title(Title)
plt.xlabel(Xl)
plt.ylabel(Yl)

# Rotation of the bars names
plt.xticks(y_pos,Top5_Stations , rotation=40, ha = 'right')
# Show graphic
plt.show()
```



## According to their website

<https://www.citibikenyc.com/pricing> (<https://www.citibikenyc.com/pricing>)

## Question 2) Trip duration by user type

According to their website <https://www.citibikenyc.com/pricing> For Annual Members pass the first 45min ride is free but if the user wants to keep, they need to pay 2.50/15min. For visitor day pass 12 for 24 hours, the 30min of each ride included in this pass but if the user wants to keep, they need to pay 2.50/15min. For 3 days pass is 24/72hrs, the 30min of each ride included in this pass but if the user wants to keep, they need to pay 4/15min. The reason for discussing all rates is here as very few people want to ride a bike more than 1 hour or safer side 2 hours because it is not economical. Also if the bike is away for more than 2 hours so it may be a chance of stolen. So it means there are lots of inconsistency in Trip duration so we are dividing our data in two portions one with abnormalities and other with normalities for visualizing the part of the case study. 1) First Half- with anomalies in dataset The graph under ax2 is a bargraph of average trip duration for each user type. It's helpful, but would be better to see a boxplot and get an idea of the distribution and see minutes instead of seconds. 2) Second graph is a basic Boxplot based with anomalies included. As we can see, there is too much noise for this to be

useful. It'll be better to look at this without anomalies. Second Half - without anomalies in dataset Still not useful, let's add a column with minutes for trip Duration. Boxplot with minutes is much more useful. There are still some outliers, however, it is informative.

### calculate Trip Duration

```
In [77]: #This question is a bit unclear in terms of what to do with the anomalies/inconsistency, so I'll be  
#making two graphs. One with anomalies, one without.  
TD_user = pd.DataFrame()  
TD_user['Avg. Trip Duration'] = round(df.groupby('User Type')['Trip Duration'].mean(),2)  
TD_user = TD_user.reset_index()  
TD_user['User Type'] = TD_user['User Type'].astype('object')
```

```
In [78]: TD_user['Avg. Trip Duration']
```

```
Out[78]: 0    2525.06  
        1     718.81  
        Name: Avg. Trip Duration, dtype: float64
```

```
In [79]: TD_user
```

```
Out[79]:
```

	User Type	Avg. Trip Duration
0	Customer	2525.06
1	Subscriber	718.81

```
In [80]: TD_user['User Type']
```

```
Out[80]: 0    Customer  
        1    Subscriber  
        Name: User Type, dtype: object
```

```
In [81]: df.groupby('User Type')['Trip Duration']
```

```
Out[81]: <pandas.core.groupby.SeriesGroupBy object at 0x000001EB5BC89898>
```

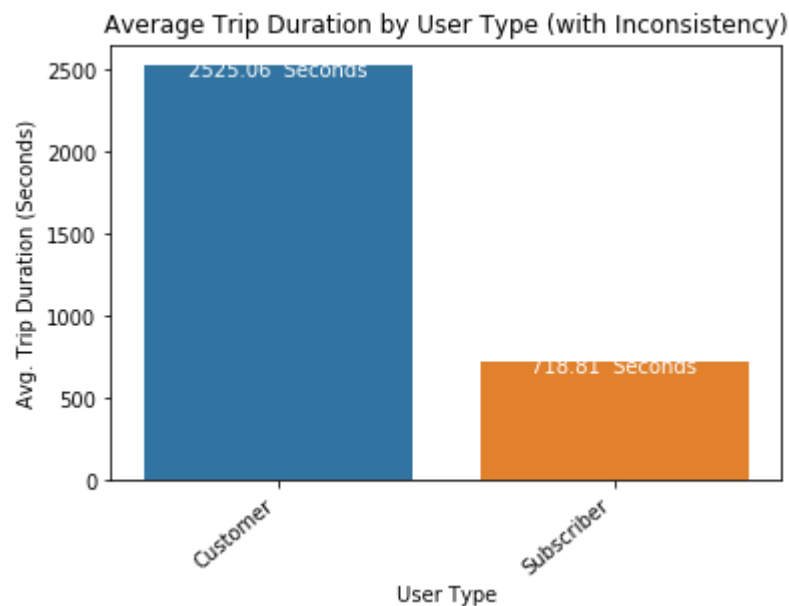
```
In [82]: df.groupby('User Type')['Trip Duration'].mean()
```

```
Out[82]: User Type
Customer      2525.061078
Subscriber     718.807176
Name: Trip Duration, dtype: float64
```

## Average trip Duration (secs)per User Type with Anomalies/Inconsistency


```
In [83]: #Average trip Duration per User Type with Anomalies/Inconsistency
ax2 = sns.barplot('User Type', 'Avg. Trip Duration', data = TD_user)
ax2.set_title('Average Trip Duration by User Type (with Inconsistency)')

ax2.set_xticklabels(ax2.get_xticklabels(),rotation=40, ha = 'right')
ax2.set_ylabel('Avg. Trip Duration (Seconds)')
for index, row in TD_user.iterrows():
    ax2.text(index,row['Avg. Trip Duration']-70,(str(row['Avg. Trip Duration'])+" Seconds"),
             color='white', ha="center", fontsize = 10)
plt.show()
```



## Boxplots are more informative to visualize breakdown of data

```
In [84]: #Boxplots are more informative to visualize breakdown of data
del(TD_user)
df.boxplot('Trip Duration', by = 'User Type')
plt.show()
```





**Any trip which lasts longer than 2 hours (7,200 seconds). Remove Inconsistency /Anomalies**

```
In [85]: df = df.drop(df.index[(df['Trip Duration'] > 7200)])
```

## Boxplots are more informative to visualize breakdown of data

```
In [86]: df.boxplot('Trip Duration', by = 'User Type')
plt.show()
#Boxplot without outliers
df.boxplot('Trip Duration', by = 'User Type', showfliers=False)
plt.show()
```



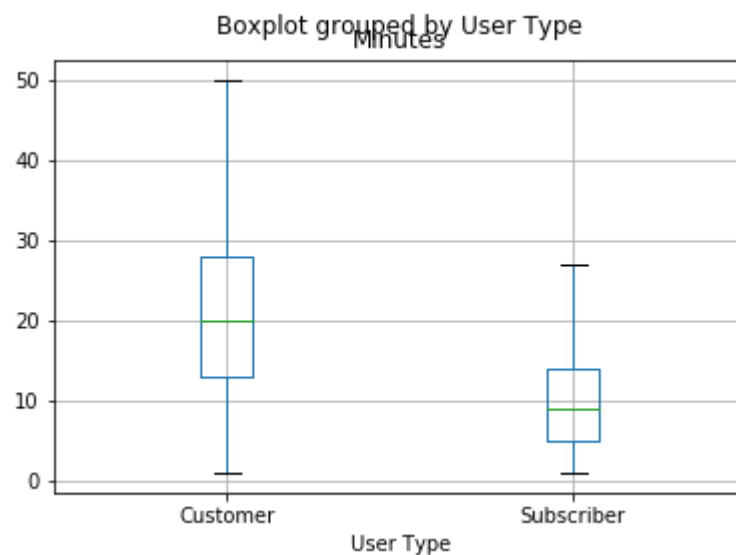

**Add Minutes column for Trip Duration**

```
In [87]: df['Minutes'] = df['Trip Duration']/60  
         #For Visual purposes, rounded  
         df['Minutes'] = round(df['Minutes'])  
         df['Minutes'] = df['Minutes'].astype(int)
```

**Final Boxplot with some outliers. Could turn of outliers with showfliers = False**



```
In [88]: df.boxplot('Minutes', by = 'User Type')  
plt.show()  
df.boxplot('Minutes', by = 'User Type', showfliers = False)  
plt.show()
```



```
In [89]: TD_user2 = pd.DataFrame()
TD_user2['Avg. Trip Duration'] = round(df.groupby('User Type')['Minutes'].mean(),1)
TD_user2 = TD_user2.reset_index()
TD_user2['User Type'] = TD_user2['User Type'].astype('object')
```

```
In [90]: TD_user2['Avg. Trip Duration']
```

```
Out[90]: 0    23.2
         1    10.9
         Name: Avg. Trip Duration, dtype: float64
```

```
In [91]: TD_user2['User Type']
```

```
Out[91]: 0    Customer
         1    Subscriber
         Name: User Type, dtype: object
```

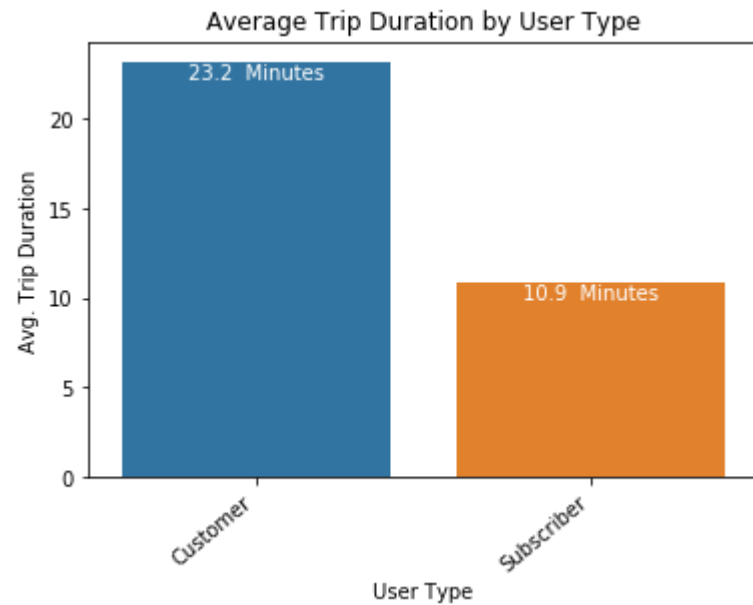
```
In [92]: TD_user2
```

```
Out[92]:
```

	User Type	Avg. Trip Duration
0	Customer	23.2
1	Subscriber	10.9

## Average Trip Duration Based on Minutes

```
In [93]: #Average Trip Duration Based on Minutes
ax3 = sns.barplot('User Type', 'Avg. Trip Duration', data = TD_user2)
ax3.set_title('Average Trip Duration by User Type')
#rcParams['figure.figsize'] = 12,10
ax3.set_xticklabels(ax3.get_xticklabels(),rotation=40, ha = 'right')
for index, row in TD_user2.iterrows():
    ax3.text(row.name,row['Avg. Trip Duration']-1,(str(row['Avg. Trip Duration'])+" Minutes"),
            color='white', ha="center", fontsize = 10)
plt.show()
```



```
In [94]: del(TD_user2)
```

## undo rounding for modelling purposes

```
In [95]: df['Minutes'] = df['Trip Duration']/60
```

## QUESTION 3: Most Popular Trip

To get most popular trips, the most convenient way to do this is by using the groupby function in pandas. It's analogous to a Pivot table.

The groupby function makes it extremely easy and convenient to identify the most popular trips. Visuals and transformations can be found below

```
In [96]: #Identify the 10 most popular trips
trips_df = pd.DataFrame()
trips_df = df.groupby(['Start Station Name', 'End Station Name']).size().reset_index(name = 'Number of Trips')
trips_df = trips_df.sort_values('Number of Trips', ascending = False)
```

```
In [97]: trips_df.head()
```

Out[97]:

	Start Station Name	End Station Name	Number of Trips
<b>46225</b>	E 7 St & Avenue A	Cooper Square & E 7 St	440
<b>83306</b>	W 21 St & 6 Ave	9 Ave & W 22 St	367
<b>55741</b>	Greenwich Ave & Charles St	Greenwich Ave & Charles St	353
<b>38414</b>	E 33 St & 2 Ave	W 33 St & 7 Ave	317
<b>70416</b>	Pershing Square North	E 24 St & Park Ave S	305

```
In [98]: trips_df["Start Station Name"] = trips_df["Start Station Name"].astype(str)
trips_df["End Station Name"] = trips_df["End Station Name"].astype(str)
```

```
In [99]: trips_df["Trip"] = trips_df["Start Station Name"] + " to " + trips_df["End Station Name"]
```

```
In [100]: trips_df = trips_df[:10]
          print(trips_df)
```

	Start Station Name	End Station Name \
46225	E 7 St & Avenue A	Cooper Square & E 7 St
83306	W 21 St & 6 Ave	9 Ave & W 22 St
55741	Greenwich Ave & Charles St	Greenwich Ave & Charles St
38414	E 33 St & 2 Ave	W 33 St & 7 Ave
70416	Pershing Square North	E 24 St & Park Ave S
68002	N 6 St & Bedford Ave	Wythe Ave & Metropolitan Ave
70348	Pershing Square North	Broadway & W 32 St
70606	Pershing Square North	W 33 St & 7 Ave
98238	Wythe Ave & Metropolitan Ave	N 6 St & Bedford Ave
72758	Richardson St & N Henry St	Graham Ave & Conselyea St

	Number of Trips	Trip
46225	440	E 7 St & Avenue A to Cooper Square & E 7 St
83306	367	W 21 St & 6 Ave to 9 Ave & W 22 St
55741	353	Greenwich Ave & Charles St to Greenwich Ave & ...
38414	317	E 33 St & 2 Ave to W 33 St & 7 Ave
70416	305	Pershing Square North to E 24 St & Park Ave S
68002	295	N 6 St & Bedford Ave to Wythe Ave & Metropolit...
70348	293	Pershing Square North to Broadway & W 32 St
70606	276	Pershing Square North to W 33 St & 7 Ave
98238	255	Wythe Ave & Metropolitan Ave to N 6 St & Bedfo...
72758	253	Richardson St & N Henry St to Graham Ave & Con...

```
In [101]: trips_df = trips_df.drop(['Start Station Name', "End Station Name"], axis = 1)
```

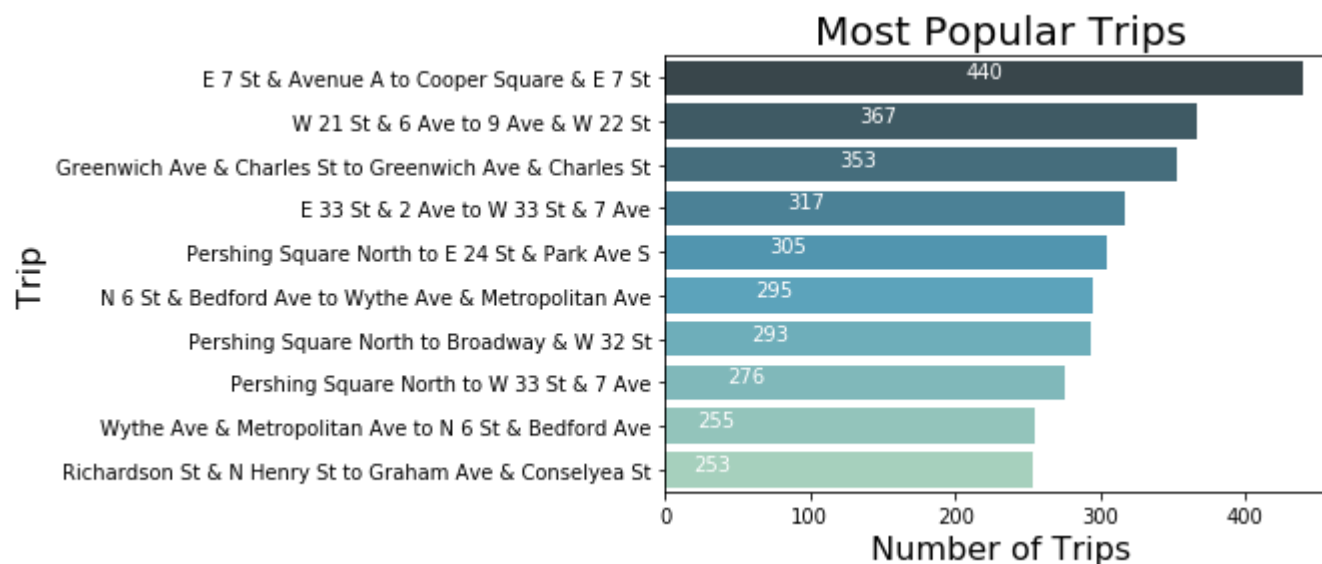
```
In [102]: trips_df = trips_df.reset_index()
```

In [103]: `trips_df.head()`

Out[103]:

	index	Number of Trips	Trip
0	46225	440	E 7 St & Avenue A to Cooper Square & E 7 St
1	83306	367	W 21 St & 6 Ave to 9 Ave & W 22 St
2	55741	353	Greenwich Ave & Charles St to Greenwich Ave & ...
3	38414	317	E 33 St & 2 Ave to W 33 St & 7 Ave
4	70416	305	Pershing Square North to E 24 St & Park Ave S

```
In [104]: ax4 = sns.barplot('Number of Trips', 'Trip', data = trips_df, palette="GnBu_d")
ax4.set_title('Most Popular Trips', fontsize = 20)
ax4.set_ylabel("Trip", fontsize=16)
ax4.set_xlabel("Number of Trips", fontsize=16)
for index, row in trips_df.iterrows():
    ax4.text(row['Number of Trips']-220, index, row['Number of Trips'],
            color='white', ha="center", fontsize = 10)
plt.show()
```



## Deleting Trip\_df dataframe

```
In [105]: del(trips_df)
```

## QUESTION 4: Rider Performance by Gender and Age

ASK;- Rider performance by Gender and Age based on avg trip distance (station to station), median speed (trip duration/distance traveled) Let's make sure the data we're working with here is clean. Ask: Rider performance by Gender and Age based on avg trip distance (station to station), median speed (trip duration/distance traveled) Let's make sure the data we're working with here is clean. Missing Gender and Birth Year values - Check missing\_table above No for Gender. Yes for Birth Year ~10% Missing Birth year. Not a big chunk of data. Can either impute missing values or drop it. Since it's less than 10% of the data, it's safe to assume the rest of the 90% is a representative sample of data and we can replace the birth year with the median, based on gender and Start Station ID. I chose this method because most people the same age live in similar neighborhoods (i.e: young people in east village, older people in Upper West Side, etc.). This will be done after anomalies are removed and speed is calculated. Are there anomalies? For Birth Year, there are some people born prior to 1957. I can believe some 60 year olds can ride a bike and that's a stretch, however, anyone "born" prior to that riding a citibike is an anomaly and false data. There could be a few senior citizens riding a bike, but probably not likely. My approach is to identify the age 2 standard deviations lower than the mean. After calculating this number, mean-2stddev, I removed the tail end of the data, birth year prior to 1957. Calculate an Age column to make visuals easier to interpret. Calculate trip distance (Miles) Calculate Speed (min/mile) and (mile/hr) (min/mile): Can be used like sprint time (how fast does this person run) (mile/hr): Conventional approach. Miles/hour is an easy to understand unit of measure and one most people are used to seeing. So the visual will be created based on this understanding. Dealing with "circular" trips Circular trips are trips which start and end at the same station. The distance for these trips will come out to 0, however, that is not the case. These points will skew the data and visuals. Will be removing them to account for this issue. For the model, this data is also irrelevant. Because if someone is going on a circular trip, the only person who knows how long the trip is going to take is themselves, assuming they know that. So it's safe to drop this data for the model.

```
In [106]: #Drop the tail end of birth years 2 standard deviations below the mean
#df['Birth Year'].mean()-(2*df['Birth Year'].std())
df = df.drop(df.index[(df['Birth Year'] < 1957)])
```

## Reset Index to avoid issues in future calculations

```
In [107]: #Reset Index to avoid issues in future calculations
df = df.reset_index()
df = df.drop('index',axis =1)
```

```
In [108]: #Combine coordinates to calculate distance based on Vincenty
df['Start Coordinates'] = list(zip(df['Start Station Latitude'], df['Start Station Longitude']))
df['End Coordinates'] = list(zip(df['End Station Latitude'], df['End Station Longitude']))
```

### Install geopy to find distance using given coordinates

```
In [109]: !pip install geopy
from geopy.distance import geodesic

dist = []
for i in range(len(df)):
    dist.append(geodesic(df['Start Coordinates'][i],df['End Coordinates'][i]).miles)
    if (i%1000000==0):
        print(i)
```

Requirement already satisfied: geopy in c:\users\mahwish\anaconda3\lib\site-packages

Requirement already satisfied: geographiclib<2,>=1.49 in c:\users\mahwish\anaconda3\lib\site-packages (from geopy)

0

You are using pip version 9.0.3, however version 18.1 is available.

You should consider upgrading via the 'python -m pip install --upgrade pip' command.

```
In [110]: #print(dist[i])
```

```
In [111]: #Reset Index to avoid issues in future calculations
df = df.reset_index()
df = df.drop('index',axis =1)
```

```
In [112]: df['Distance'] = dist
```

```
In [113]: del(dist)
#Replace missing birth year by median based on speed and gender
df['Birth Year'] = df.groupby(['Gender','Start Station ID'])['Birth Year'].transform(lambda x: x.fillna(x.median()))
```



```
In [114]: #Still have a few nulls, but few so Comfortable dropping these.  
df = df.dropna(subset=['Birth Year'])
```

```
In [115]: #Calculate age and drop circular/roundtrips  
df['Age'] = 2018 - df['Birth Year']  
df['Age'] = df['Age'].astype(int)
```

```
In [116]: df = df.drop(df.index[(df['Distance'] == 0)])
```

**. Followed the same reasoning as behind Birth Year. People in similar locations tend to also work in a similar industry or location**

```
In [117]: df['Distance'] = df.groupby(['Gender', 'Start Station ID'])['Distance'].transform(lambda x: x.fillna(x.median()))
```

### **Caulculate Speed (min/mile) and (mile/hr)**

(min/mile): Can be used like sprint time (how fast does this person run) (mile/hr): Conventional approach. Miles/hour is an easy to understand unit of measure and one most people are used to seeing. So the visual will be created based on this understanding.

```
In [118]: df['min_mile'] = round(df['Minutes']/df['Distance'], 2)  
df['mile_hour'] = round(df['Distance']/(df['Minutes']/60),2)
```

**Let's check for data integrity to make sure all the numbers look as expected. Only numerical data included**

```
In [119]: #Let's check for data integrity to make sure all the numbers look as expected. Only numerical data  
#included  
print(round(df.describe(),2))
```

	Trip Duration	Start Station ID	Start Station Latitude \
count	670501.00	670501.00	670501.00
mean	664.19	1217.57	40.74
std	505.29	1274.98	0.03
min	61.00	72.00	40.65
25%	328.00	358.00	40.72
50%	518.00	482.00	40.74
75%	840.00	3090.00	40.76
max	7193.00	3443.00	40.80

	Start Station Longitude	End Station ID	End Station Latitude \
count	670501.00	670501.00	670501.00
mean	-73.98	1192.23	40.74
std	0.02	1263.49	0.03
min	-74.02	72.00	40.65
25%	-74.00	357.00	40.72
50%	-73.99	479.00	40.74
75%	-73.98	3076.00	40.75
max	-73.93	3447.00	40.80

	End Station Longitude	Bike ID	Birth Year	Gender	Minutes \
count	670501.00	670501.00	670501.00	670501.00	670501.00
mean	-73.99	21738.08	1978.65	1.17	11.07
std	0.02	4199.95	10.28	0.47	8.42
min	-74.03	14529.00	1957.00	0.00	1.02
25%	-74.00	17878.00	1970.00	1.00	5.47
50%	-73.99	21329.00	1980.00	1.00	8.63
75%	-73.98	25812.00	1987.00	1.00	14.00
max	-73.93	27325.00	2000.00	2.00	119.88

	Distance	Age	min_mile	mile_hour
count	670501.00	670501.00	670501.00	670501.00
mean	1.05	39.35	11.82	6.02
std	0.78	10.28	14.40	1.71
min	0.02	18.00	3.44	0.01
25%	0.53	31.00	8.45	5.00
50%	0.83	38.00	9.96	6.02
75%	1.31	48.00	12.01	7.10
max	8.98	61.00	4116.27	17.46

```
In [120]: df = df.drop(df.index[(df['Distance'] == 0)])
```

Observations We still have trips less than 90 seconds, however they seem to be legitimate trips. Checked using the code in cell above. We have some Start Coordinates as (0.0,0.0). These are trips which were taken away for repair or for other purposes. These should be dropped. If kept, the distance for these trips is 5,389 miles. For this reason I've dropped any points where the distance is greater than 30 miles. Additionally, we have some missing values. Since it's a tiny portion, let's replace missing values based on Gender and start location. These One some trips, the speed of the biker is more than 200 mph. This could be due to the formula used for distance calculation or some other error. The fastest cyclist in the world on a flat surface ever recorded biked at 82mph. It's safe to assume none of the citibike riders can approach this speed. Due to this and the fact that an average cyclist speed is 10mph, I've decided to remove all data where the speed in mph is greater than 20 mph and less than 0.1 mph. ~1.5k data points


```
In [121]: df = df[df['Distance'] < 30]
```

```
In [122]: #1-Done in two steps to ensure data integrity, could've used an or statement as well.  
df = df[df['mile_hour']<20]  
#2  
df = df[df['mile_hour']> (df['mile_hour'].mean()-(2*df['mile_hour'].std()))]
```


```
In [123]: #Dropping unknown to make the visual more informative.  
#Unknown gender may be important for the model, which is why I created a copy of the original dataframe.  
df1 = df.drop(df.index[(df['Gender'] == 0)])
```

```
In [124]: #Rider performance by age and Gender in Min/Mile
```


```
In [125]: #Rider performance by age and Gender in Min/Mile
fig, ax5 = plt.subplots(figsize=(11,5))
df1.groupby(['Age', 'Gender']).median()['min_mile'].unstack().plot(ax=ax5, color = "bg")
ax5.legend(['Female', 'Male'])
plt.ylabel('Median Speed (min/mile)')
plt.title('Rider Performance Based on Gender and Age (Median Speed in min/mile)')
plt.show()
```



```
In [126]: #Rider performance by age and Gender in Miles/hr
del([fig,ax5])
fig1, ax6 = plt.subplots(figsize=(11,5))
df1.groupby(['Age','Gender']).median()['mile_hour'].unstack().plot(ax=ax6,color="bg")
ax6.legend(['Female','Male'])
plt.ylabel('Median Speed (miles/hr)')
plt.title('Rider Performance Based on Gender and Age (Median Speed in miles/hr)')
plt.show()
```



```
In [127]: #Rider performance by age and Gender in Average Distance
del([fig1,ax6])
fig2, ax7 = plt.subplots(figsize=(11,5))
df1.groupby(['Age', 'Gender']).mean()['Distance'].unstack().plot(ax=ax7,color = "bg")
ax7.legend(['Female', 'Male'])
plt.ylabel('Average Distance (miles)')
plt.title('Rider Performance Based on Gender and Age (Average Distance in Miles)')
plt.show()
```



## QUESTION 5

### Busiest Bike by Times and Minutes Used

## What is the busiest bike in NYC in 2017?

Bike 26022

How many times was it used? times 250

How many minutes was it in use? Minutes 2984

## Busiest bike and count can be identified by a groupby function

Function above will also identify the number of times the bike was used A similar groupby function which calls for the sum on minutes can identify the number of minutes the bike was used.

### Bike usage based on number of times used

```
In [128]: #Bike usage based on number of times used
del(df1)
bike_use_df = pd.DataFrame()
bike_use_df = df.groupby(['Bike ID']).size().reset_index(name = 'Number of Times Used')
bike_use_df = bike_use_df.sort_values('Number of Times Used', ascending = False)
#bike_use_df.to_csv('Q5.csv')
bike_use_df = bike_use_df[:10]
bike_use_df['Bike ID'] = bike_use_df['Bike ID'].astype(str)
bike_use_df['Bike ID'] = ('Bike ' + bike_use_df['Bike ID'])
bike_use_df = bike_use_df.reset_index()
```




```
In [129]: bike_use_df.head()
```

```
Out[129]:
```

	index	Bike ID	Number of Times Used
0	7088	Bike 26022	250
1	7470	Bike 26585	243
2	7282	Bike 26386	243
3	7460	Bike 26573	241
4	7221	Bike 26323	238

### Visual of most used bike based on Number of Trips

```
In [130]: ax8 = sns.barplot('Number of Times Used', 'Bike ID', data = bike_use_df, palette="GnBu_d")
ax8.set_title('Most Popular Bikes by Number of Times Used')
for index, row in bike_use_df.iterrows():
    ax8.text(row['Number of Times Used']-90, index, row['Number of Times Used'], color='white', ha="center", fo
ntsize =10)
plt.show()
```



**Bike usage based on minutes used**

In [131]: `del(ax8)`

```
In [132]: bike_min_df = pd.DataFrame()
bike_min_df['Minutes Used'] = df.groupby('Bike ID')['Minutes'].sum()
bike_min_df = bike_min_df.reset_index()
bike_min_df = bike_min_df.sort_values('Minutes Used', ascending = False)
bike_min_df['Bike ID'] = bike_min_df['Bike ID'].astype(str)
bike_min_df['Bike ID'] = ('Bike ' + bike_min_df['Bike ID'])
bike_min_df = bike_min_df[:10]
bike_min_df = bike_min_df.reset_index()
```

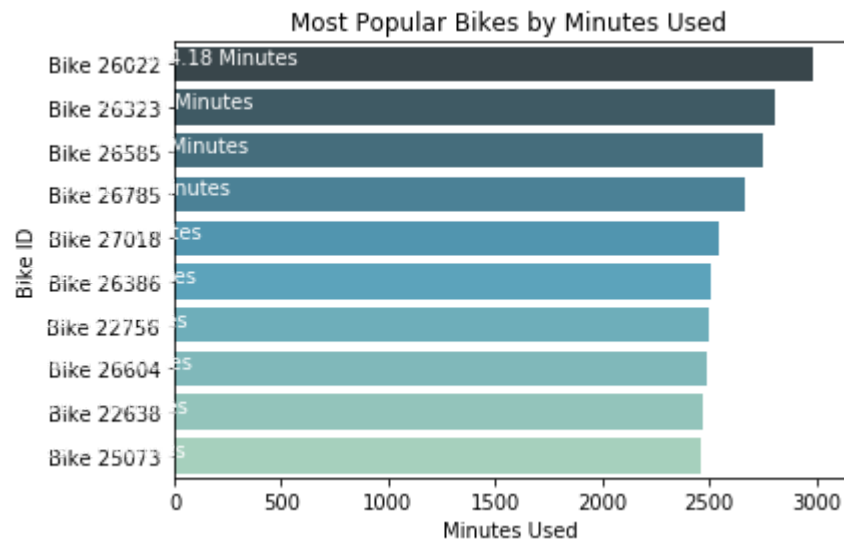
In [133]: `bike_min_df.head()`

Out[133]:

	index	Bike ID	Minutes Used
0	7088	Bike 26022	2984.183333
1	7221	Bike 26323	2810.500000
2	7470	Bike 26585	2751.966667
3	7648	Bike 26785	2671.283333
4	7878	Bike 27018	2540.733333

**#Visual of most used bike based on number of minutes used**

```
In [134]: ax9 = sns.barplot('Minutes Used', 'Bike ID', data = bike_min_df, palette="GnBu_d")
ax9.set_title('Most Popular Bikes by Minutes Used')
rcParams['figure.figsize'] = 11,6
for index, row in bike_min_df.iterrows():
    ax9.text(row['Minutes Used']-2800, index, str(round(row['Minutes Used'],2))+ ' Minutes',
            color='white', ha="center")
plt.show()
```



## 1) Predictive Model - Baseline Model

**Ask: Build a model that can predict how long a trip will take given a starting point and destination.**

Let's now begin to train our regression model! We will need to first split up our data into an X array that contains the features to train on, and a y array with the target variable, in this case the Minutes column. Assumptions on how the Kiosk will work: After speaking to Daniel Yawitz (if you're looking at this, thanks for the clarification), I was told that we should assume that when a user inputs the start and end station, they swipe their key fob (if they're a subscriber) and enter their info on the kiosk (if they're a "Customer") prior to entering the start and end station. This means that we would know their gender and age. Thus these variables can be used in building the model.

**Step 1.**

This dataset is massive, Let's work on a random subsample while we build and evaluate models. If I tried to build and evaluate a model on the entire dataset, each run would take me ~10+ minutes depending on the model. I also made the same visuals on the sample to visualize the data.

**Step 2.**

Let's get a baseline. If I were to just run a simple multi-variate linear regression, what would my model look like and how accurate would it be? Need to prepare the data for a multivariate regression

- 1) Drop irrelevant columns
  - Trip Duration:** We have the minutes column, which is the target variable
  - Stop Time:** In the real world, we won't have this information when predicting the trip duration.
  - Start Station ID:** Start Station Name captures this information
  - Start Station Latitude:** Start Station Name captures this information
  - Start Station Longitude:** Start Station Name captures this information
  - Start Coordinates:** Start Station Name captures this information
  - End Station ID:** End Station Name captures this information
  - End Station Latitude:** End Station Name captures this information
  - End Station Longitude:** End Station Name captures this information
  - End Coordinates:** End Station Name captures this information
  - Bike Id:** We won't know what bike the user is going to end up using
  - Min\_Mile:** Effectively the same information as end time when combined with distance. We won't have this information in the real world.
  - mile\_hour:** Effectively the same information as end time when combined with distance. We won't have this information in the real world. ( $\text{Speed} * \text{Distance} = \text{Trip Duration}$ ): Which is why speed is dropped
  - Birth Year:** Age captures this information
  - Start Station Name and End Station Name:** The distance variable captures the same information. For the model, if a user is inputting start and end station, we can build a simple function to calculate the distance which would capture the same information.
- 2) Basic cleaning of data FOR NOW. This is only being done for the baseline model
  - Start Time:** Requires reformatting. Will do this after baseline model
  - Dumify categorical variables**
  - Scale Age**
  - Don't scale distance,** since it does not just represent distance, but is also indicative of the trip the rider is making (start and station)
- 3) Anomalies in Trip Duration I'm going to come back to an observation from earlier. Any trip which lasts longer than 45 minutes (2,700 seconds) probably indicates a stolen bike, incorrect docking of the bike, or an anomaly. No rider would plan to go over the maximum 45 minutes allowed. Age was removed after an initial run indicated it had no effect on the model. This was also clearly indicated in some of the visuals above.

```
In [136]: print(round(df.describe(),2))
```

	Trip Duration	Start Station ID	Start Station Latitude \
count	649110.00	649110.00	649110.00
mean	635.93	1214.20	40.74
std	451.03	1273.17	0.03
min	61.00	72.00	40.65
25%	323.00	359.00	40.72
50%	507.00	482.00	40.74
75%	810.00	3090.00	40.76
max	7062.00	3443.00	40.80

	Start Station Longitude	End Station ID	End Station Latitude \
count	649110.00	649110.00	649110.00
mean	-73.98	1188.18	40.74
std	0.02	1261.28	0.03
min	-74.02	72.00	40.65
25%	-74.00	357.00	40.72
50%	-73.99	479.00	40.74
75%	-73.98	3071.00	40.75
max	-73.93	3447.00	40.80

	End Station Longitude	Bike ID	Birth Year	Gender	Minutes \
count	649110.00	649110.00	649110.00	649110.00	649110.00
mean	-73.99	21746.11	1978.69	1.18	10.60
std	0.02	4200.11	10.26	0.46	7.52
min	-74.03	14529.00	1957.00	0.00	1.02
25%	-74.00	17886.00	1971.00	1.00	5.38
50%	-73.99	21339.00	1980.00	1.00	8.45
75%	-73.98	25816.00	1987.00	1.00	13.50
max	-73.93	27325.00	2000.00	2.00	117.70

	Distance	Age	min_mile	mile_hour
count	649110.00	649110.00	649110.00	649110.00
mean	1.06	39.31	10.40	6.17
std	0.78	10.26	2.88	1.53
min	0.06	18.00	3.44	2.60
25%	0.54	31.00	8.40	5.11
50%	0.84	38.00	9.86	6.09
75%	1.33	47.00	11.74	7.14
max	8.98	61.00	23.12	17.46

## Observations

We have some unreasonable speeds in min\_mile column, but that's fine. Some people may have walked with their bike or stopped at multiple destinations before docking. The sample data below seems to be representative of the entire dataset above. #Training a Linear Regression Model #We have the minutes column, which is the target variable #1min =60 sec #45min =45\*60= 2700 sec=Trip duration more than this number will be dropped from the data set.

## Cleaning up anomalies

```
In [137]: #Cleaning up anomalies  
df = df.drop(df.index[(df['Trip Duration'] > 2700)])
```

Training a Linear Regression Model We have the minutes column, which is the target variable 1min =60 sec 45min =45\*60= 2700 sec=Trip duration more than this number will be dropped from the data set.

```
In [142]: #Let's work with a random sample of the data and inspect it thoroughly to ensure it's representative  
df_sample = df.sample(frac = 0.1, random_state = 0)  
print(df_sample.describe())
```

	Trip Duration	Start Station ID	Start Station Latitude \
count	64764.000000	64764.000000	64764.000000
mean	629.120453	1211.78539	40.737337
std	434.848317	1273.24560	0.026285
min	61.000000	72.000000	40.646768
25%	322.000000	358.000000	40.720874
50%	505.500000	481.500000	40.739355
75%	805.000000	3088.000000	40.755003
max	2700.000000	3443.000000	40.804213

	Start Station Longitude	End Station ID	End Station Latitude \
count	64764.000000	64764.000000	64764.000000
mean	-73.985055	1189.697455	40.737094
std	0.015974	1261.254191	0.025941
min	-74.017134	72.000000	40.646768
25%	-73.995299	357.000000	40.720828
50%	-73.987216	480.000000	40.739355
75%	-73.976806	3071.000000	40.754666
max	-73.929891	3443.000000	40.804213

	End Station Longitude	Bike ID	Birth Year	Gender \
count	64764.000000	64764.000000	64764.000000	64764.000000
mean	-73.985361	21741.058860	1978.714764	1.179220
std	0.016033	4200.926265	10.247366	0.458099
min	-74.017134	14529.000000	1957.000000	0.000000
25%	-73.995960	17896.000000	1971.000000	1.000000
50%	-73.987586	21347.000000	1981.000000	1.000000
75%	-73.977061	25805.000000	1987.000000	1.000000
max	-73.929891	27325.000000	2000.000000	2.000000

	Minutes	Distance	Age	min_mile	mile_hour
count	64764.000000	64764.000000	64764.000000	64764.000000	64764.000000
mean	10.485341	1.051823	39.284031	10.399077	6.166821
std	7.247472	0.766826	10.247372	2.882319	1.533457
min	1.016667	0.067453	18.000000	4.240000	2.600000
25%	5.366667	0.533514	31.000000	8.420000	5.110000
50%	8.425000	0.835016	37.000000	9.860000	6.080000
75%	13.416667	1.320865	47.000000	11.740000	7.130000
max	45.000000	7.021076	61.000000	23.120000	14.140000



## Build a Baseline Model

```
In [143]: #Drop Irrelevant data
def drop_data(df):
    df = df.drop(['Trip Duration', 'Stop Time', 'Start Station ID', 'Start Station Latitude', 'Start Station Longitude',
                  'Start Coordinates', 'End Station ID', 'End Station Latitude', 'End Station Longitude',
                  'End Coordinates', 'Bike ID', 'Start Station Name', 'Birth Year', 'End Station Name', 'min_mile',
                  'mile_hour', 'Age'], axis = 1)
    return df

df_basemodel = drop_data(df_sample)
```

```
In [144]: df_basemodel = df_basemodel.drop('Start Time', axis =1)
```

```
In [145]: #Dummify categorical data and avoid dummy variable trap
df_basemodel = pd.get_dummies(df_basemodel, drop_first = True)
df_basemodel.dtypes
```

```
Out[145]: Gender                int64
Minutes                float64
Distance                float64
User Type_Subscriber    uint8
dtype: object
```

```
In [146]: df_basemodel.head()
```

```
Out[146]:
```

	Gender	Minutes	Distance	User Type_Subscriber
<b>259203</b>	1	8.550000	0.756068	1
<b>360096</b>	1	6.633333	0.697781	1
<b>434602</b>	1	5.900000	1.023749	1
<b>478347</b>	1	4.000000	0.370099	1
<b>347879</b>	2	15.583333	1.488621	1

```
In [147]: ##As We have the minutes column, which is the target variable
df_basemodel.corr().loc[:, 'Minutes']
```

```
Out[147]: Gender                0.005635
Minutes                1.000000
Distance              0.899097
User Type_Subscriber  -0.152969
Name: Minutes, dtype: float64
```

## Training a Linear Regression Base Model¶

Let's now begin to train our regression model! We will need to first split up our data into an X array that contains the features to train on, and a y array with the target variable, in this case the Minutes column

```
In [151]: #Train Test Split
#Predictor variable
X = df_basemodel[['Gender', 'Distance', 'User Type_Subscriber']]
#Target variable
y = df_basemodel['Minutes']
```

*#Train Test Split* *#Our goal is to create a model that generalises well to new data. Our test set serves as a proxy for new data. Trained data is the data on which we apply the linear regression algorithm. #And finally we test that algorithm on the test data. The code for splitting is as follows: #from*  
*sklearn.cross\_validation import train\_test\_split*

```
In [153]: from sklearn.cross_validation import train_test_split
```

```
C:\Users\Mahwish\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module
was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and
functions are moved. Also note that the interface of the new CV iterators are different from that of this
module. This module will be removed in 0.20.
  "This module will be removed in 0.20.", DeprecationWarning)
```

```
In [154]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
```

In [155]: *#Creating and Training the Model*

```
from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(X_train,y_train)
#The above code fits the linear regression model on the training data.
```

Out[155]: LinearRegression(copy\_X=True, fit\_intercept=True, n\_jobs=1, normalize=False)


In [156]: *#Predictions from our Model*

```
#Let's grab predictions off the test set and see how well it did!
predictions = lm.predict(X_test)
```

In [157]: *#Let's visualise the prediction*

```
plt.scatter(y_test,predictions)
```

Out[157]: <matplotlib.collections.PathCollection at 0x1eb09215160>



```
In [158]: #Model Accuracy  
lm.score(X_test,y_test)
```

```
Out[158]: 0.8226371686192211
```

```
In [159]: #Using Statsmodel because it has the summary function.
import statsmodels.api as sm
X_train = sm.add_constant(X_train)
X_test = sm.add_constant(X_test)
lm_OLS = sm.OLS(y_train, X_train).fit()
lm_OLS.summary()
```

Out[159]: OLS Regression Results

<b>Dep. Variable:</b>	Minutes	<b>R-squared:</b>	0.821
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.821
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	7.946e+04
<b>Date:</b>	Tue, 20 Nov 2018	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	23:34:00	<b>Log-Likelihood:</b>	-1.3167e+05
<b>No. Observations:</b>	51811	<b>AIC:</b>	2.634e+05
<b>Df Residuals:</b>	51807	<b>BIC:</b>	2.634e+05
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>[0.025</b>	<b>0.975]</b>
<b>const</b>	6.6541	0.098	68.047	0.000	6.462	6.846
<b>Gender</b>	1.0105	0.032	31.987	0.000	0.949	1.072
<b>Distance</b>	8.4293	0.018	479.888	0.000	8.395	8.464
<b>User Type_Subscriber</b>	-6.3531	0.103	-61.659	0.000	-6.555	-6.151

<b>Omnibus:</b>	16487.274	<b>Durbin-Watson:</b>	2.010
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	133659.054
<b>Skew:</b>	1.306	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	10.423	<b>Cond. No.</b>	22.3

Hence the model is successfully built